

# Introduction to DGE

[View on GitHub](#)

Approximate time: 120 minutes

## Learning Objectives:


- Determine how functions are attributed to genes using Gene Ontology terms
- Understand the theory of how functional enrichment tools yield statistically enriched functions or interactions
- Discuss functional analysis using over-representation analysis, functional class scoring, and pathway topology methods
- Explore functional analysis tools for over-representation analysis

## Functional analysis

The output of RNA-seq differential expression analysis is a list of significant differentially expressed genes (DEGs). To gain greater biological insight on the differentially expressed genes there are various analyses that can be done:

- determine whether there is enrichment of known biological functions, interactions, or pathways
- identify genes' involvement in novel pathways or networks by grouping genes together based on similar trends
- use global changes in gene expression by visualizing all genes being significantly up- or down-regulated in the context of external interaction data

Generally for any differential expression analysis, it is useful to interpret the resulting gene lists using freely available web- and R-based tools. While tools for functional analysis span a wide variety of techniques, they can loosely be categorized into three main types: over-representation analysis, functional class scoring, and pathway topology [1].



The goal of functional analysis is provide biological insight, so it's necessary to analyze our results in the context of our experimental hypothesis: **FMRP and MOV10 associate and regulate the translation of a subset of RNAs**. Therefore, based on the authors' hypothesis, we may expect the enrichment of processes/pathways related to **translation, splicing, and the regulation of mRNAs**, which we would need to validate experimentally.

*Note that all tools described below are great tools to validate experimental results and to make hypotheses. These tools suggest genes/pathways that may be involved with your condition of interest; however, you should NOT use these tools to make conclusions about the pathways involved in your experimental process. You will need to perform experimental validation of any suggested pathways.*



## Over-representation analysis

There are a plethora of functional enrichment tools that perform some type of “over-representation” analysis by querying databases containing information about gene function and interactions.

These databases typically **categorize genes into groups (gene sets)** based on shared function, or involvement in a pathway, or presence in a specific cellular location, or other categorizations, e.g. functional pathways, etc. Essentially, known genes are binned into categories that have been consistently named (controlled vocabulary) based on how the gene has been annotated functionally. These categories are independent of any organism, however each organism has distinct categorizations available.

To determine whether any categories are over-represented, you can determine the probability of having the observed proportion of genes associated with a specific category in your gene list based on the proportion of genes associated with the same category in the background set (gene categorizations for the appropriate organism).

All known genes in a species  
(categorized into groups)



DEGs

Genes categories	Organism-specific Background	DE results	Over-represented?
Functional category 1	35/13000	25/1000	Likely
Functional category 2	56/13000	4/1000	Unlikely
Functional category 3	90/13000	8/1000	Unlikely
Functional category 4	15/13000	10/1000	Likely
...			
...			

The statistical test that will determine whether something is actually over-represented is the *Hypergeometric test*.

Hypergeometric testing

Using the example of the first functional category above, hypergeometric distribution is a probability distribution that describes the probability of 25 genes (k) being associated with “Functional category 1”, for all genes in our gene list (n=1000), from a population of all of the genes in entire genome (N=13,000) which contains 35 genes (K) associated with “Functional category 1” [4].

The calculation of probability of k successes follows the formula:

$$P(X = k) = \frac{\binom{K}{k} \binom{N-K}{n-k}}{\binom{N}{n}}$$

This test will result in an adjusted p-value (after multiple test correction) for each category tested.

## Gene Ontology project

One of the most widely-used categorizations is the **Gene Ontology (GO)** established by the Gene Ontology project.

“The Gene Ontology project is a collaborative effort to address the need for consistent descriptions of gene products across databases” [2]. The [Gene Ontology Consortium](#) maintains the GO terms, and these GO terms are incorporated into gene annotations in many of the popular repositories for animal, plant, and microbial genomes.

Tools that investigate enrichment of biological functions or interactions often use the Gene Ontology (GO) categorizations, i.e. the GO terms to determine whether any have significantly modified representation in a given list of genes. Therefore, to best use and interpret the results from these functional analysis tools, it is helpful to have a good understanding of the GO terms themselves and their organization.

## GO Ontologies

To describe the roles of genes and gene products, GO terms are organized into three independent controlled vocabularies (ontologies) in a species-independent manner:

- **Biological process:** refers to the biological role involving the gene or gene product, and could include “transcription”, “signal transduction”, and “apoptosis”. A biological process generally involves a chemical or physical change of the starting material or input.
- **Molecular function:** represents the biochemical activity of the gene product, such activities could include “ligand”, “GTPase”, and “transporter”.
- **Cellular component:** refers to the location in the cell of the gene product. Cellular components could include “nucleus”, “lysosome”, and “plasma membrane”.

Each GO term has a term name (e.g. **DNA repair**) and a unique term accession number (**GO:0005125**), and a single gene product can be associated with many GO terms, since a single gene product “may function in several processes, contain domains that carry out diverse molecular functions, and participate in multiple alternative interactions with other proteins, organelles or locations in the cell” [3].

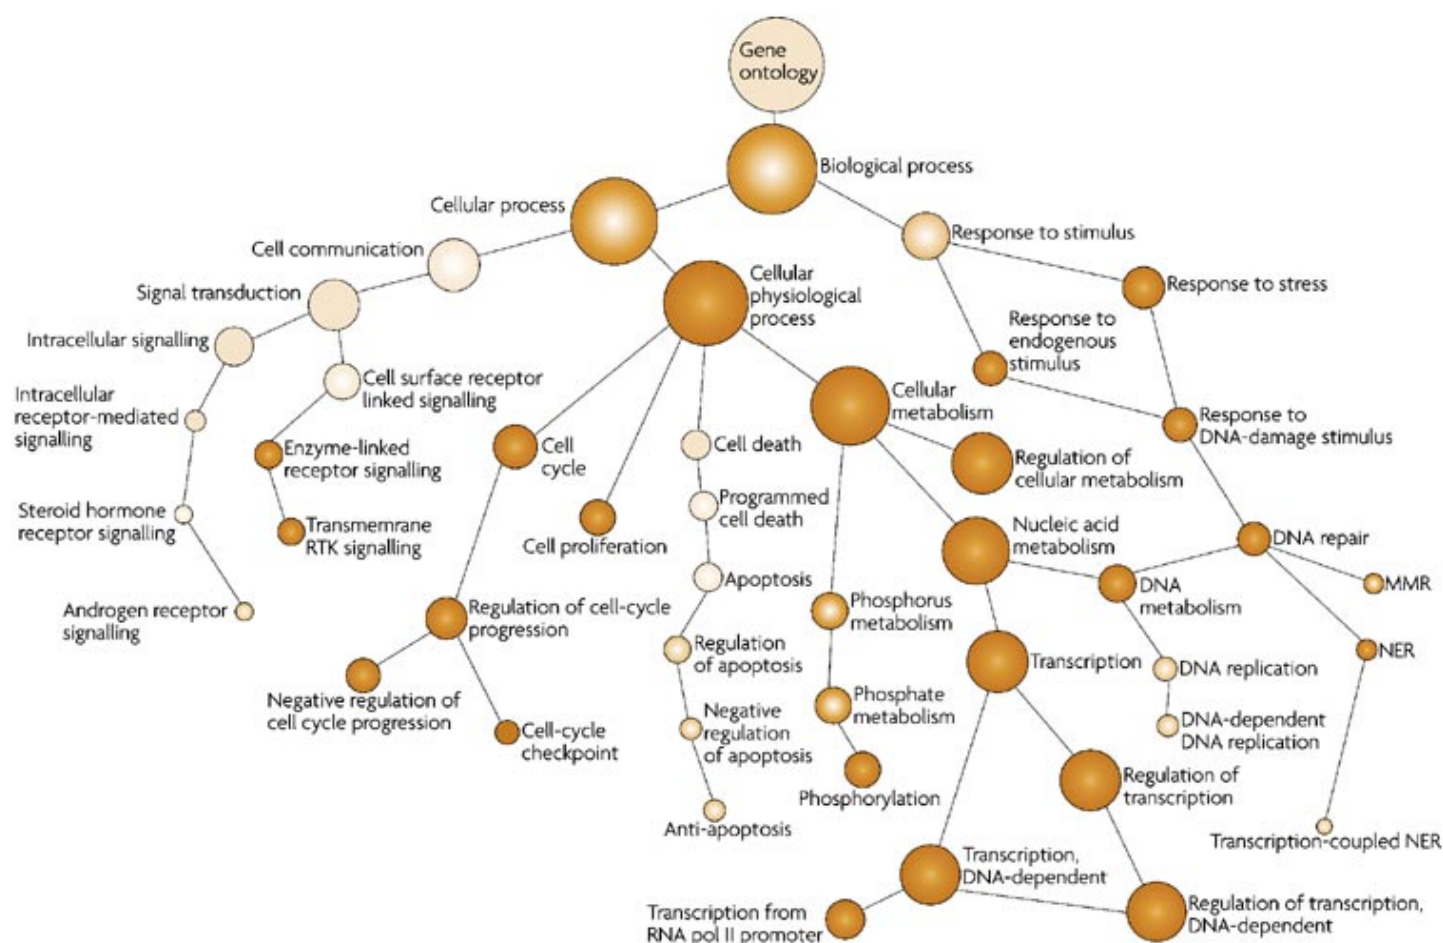
## GO term hierarchy

Some gene products are well-researched, with vast quantities of data available regarding their biological processes and functions. However, other gene products have very little data available about their roles in the cell.

For example, the protein, “p53”, would contain a wealth of information on it’s roles in the cell, whereas another protein might only be known as a “membrane-bound protein” with no other information available.

The GO ontologies were developed to describe and query biological knowledge with differing levels of information available. To do this, GO ontologies are loosely hierarchical, ranging from general, ‘parent’, terms to more specific, ‘child’ terms. The GO ontologies are “loosely” hierarchical since ‘child’ terms can have multiple ‘parent’ terms.

Some genes with less information may only be associated with general ‘parent’ terms or no terms at all, while other genes with a lot of information be associated with many terms.



Nature Reviews | Cancer

## Tips for working with GO terms

## clusterProfiler

We will be using [clusterProfiler](#) to perform over-representation analysis on GO terms associated with our list of significant genes. The tool takes as input a significant gene list and a background

gene list and performs statistical enrichment analysis using hypergeometric testing. The basic arguments allow the user to select the appropriate organism and GO ontology (BP, CC, MF) to test.

## Running clusterProfiler

To run clusterProfiler GO over-representation analysis, we will change our gene names into Ensembl IDs, since the tool works a bit easier with the Ensembl IDs.

Then load the following libraries:

```
# Load libraries
library(DOSE)
library(pathview)
library(clusterProfiler)
library(org.Hs.eg.db)
```

For the different steps in the functional analysis, we require Ensembl and Entrez IDs. We will use the gene annotations that we generated previously to merge with our differential expression results.

```
## Merge the AnnotationHub dataframe with the results
res_ids <- left_join(res_table0E_tb, annotations_ahb, by=c("gene"="gene_id"))
```

**NOTE:** If you were unable to generate the `annotations_ahb` object, you can download the annotations to your data folder by right-clicking [here](#) and selecting "Save link as..."

To read in the object, you can run the following code: `annotations_ahb <- read.csv("annotations_ahb.csv")`

To perform the over-representation analysis, we need a list of background genes and a list of significant genes. For our background dataset we will use all genes tested for differential expression (all genes in our results table). For our significant gene list we will use genes with p-adjusted values less than 0.05 (we could include a fold change threshold too if we have many DE genes).

```
## Create background dataset for hypergeometric testing using all genes tested for
all0E_genes <- as.character(res_ids$gene)

## Extract significant results
sig0E <- dplyr::filter(res_ids, padj < 0.05)

sig0E_genes <- as.character(sig0E$gene)
```



Now we can perform the GO enrichment analysis and save the results:

```
## Run GO enrichment analysis
ego <- enrichGO(gene = sigOE_genes,
  universe = allOE_genes,
  keyType = "ENSEMBL",
  OrgDb = org.Hs.eg.db,
  ont = "BP",
  pAdjustMethod = "BH",
  qvalueCutoff = 0.05,
  readable = TRUE)
```

**NOTE:** The different organisms with annotation databases available to use with for the `OrgDb` argument can be found [here](#).

Also, the `keyType` argument may be coded as `keytype` in different versions of `clusterProfiler`.

Finally, the `ont` argument can accept either “BP” (Biological Process), “MF” (Molecular Function), and “CC” (Cellular Component) subontologies, or “ALL” for all three.

```
## Output results from GO analysis to a table
cluster_summary <- data.frame(ego)

write.csv(cluster_summary, "results/clusterProfiler_Mov10oe.csv")
```

	ID	Description	GeneRatio	BgRatio	pvalue	p.adjust	qvalue	geneID	Count
GO:0008380	GO:0008380	RNA splicing	217/5660	393/16649	2.299032e-18	1.015934e-14	8.427986e-15	RBM11/RBM15B/RBM38/SNRPD1/SNRPD3/PP...	217
GO:0006397	GO:0006397	mRNA processing	247/5660	463/16649	3.630282e-18	1.015934e-14	8.427986e-15	RBM11/RBM15B/APP/RBM38/SNRPD1/FASTK...	247
GO:0010608	GO:0010608	posttranscriptional regulation of gene expres...	247/5660	481/16649	1.544298e-15	2.881145e-12	2.390140e-12	RPS27L/SMAD2/APP/RBM38/PSMB7/LSM14A...	247
GO:0034660	GO:0034660	ncRNA metabolic process	237/5660	473/16649	1.884052e-13	2.636260e-10	2.186988e-10	SMAD2/RAE1/SNRPD1/SMARCB1/RRP7BP/RP...	237
GO:0000375	GO:0000375	RNA splicing, via transesterification reactions	158/5660	292/16649	9.323903e-13	9.270219e-10	7.690386e-10	RBM11/RBM15B/SNRPD1/SNRPD3/SNRPC/LS...	158
GO:0000377	GO:0000377	RNA splicing, via transesterification reactions...	156/5660	288/16649	1.159398e-12	9.270219e-10	7.690386e-10	RBM11/RBM15B/SNRPD1/SNRPD3/SNRPC/LS...	156
GO:0000398	GO:0000398	mRNA splicing, via spliceosome	156/5660	288/16649	1.159398e-12	9.270219e-10	7.690386e-10	RBM11/RBM15B/SNRPD1/SNRPD3/SNRPC/LS...	156
GO:0022613	GO:0022613	ribonucleoprotein complex biogenesis	194/5660	379/16649	2.554931e-12	1.787494e-09	1.482869e-09	RPS27L/LSM14A/SNRPD1/RRP7BP/RPSA/WDR...	194
GO:0044772	GO:0044772	mitotic cell cycle phase transition	241/5660	499/16649	1.576324e-11	9.802982e-09	8.132356e-09	RPS27L/UBE2C/APP/INO80/RBM38/PSMB7/S...	241
GO:0018205	GO:0018205	peptidyl-lysine modification	203/5660	411/16649	5.414440e-11	3.030462e-08	2.514010e-08	RAE1/CTCF/SMARCB1/RTF1/BAC6/EEF1AK...	203
GO:0002486	GO:0002486	antigen processing and presentation of endo...	21/5660	21/16649	1.410396e-10	7.176352e-08	5.953356e-08	HLA-C/HLA-A/HLA-B/HLA-A/HLA-C/HLA-A/HL...	21
GO:0034470	GO:0034470	ncRNA processing	163/5660	320/16649	2.315204e-10	1.079850e-07	8.958215e-08	SMAD2/RAE1/RRP7BP/RPSA/WDR46/FBL/NUP...	163
GO:0016570	GO:0016570	histone modification	201/5660	412/16649	2.625099e-10	1.130206e-07	9.375959e-08	CTCF/DAXX/SMARCB1/RTF1/ZNF335/ATG7...	201
GO:0006281	GO:0006281	DNA repair	235/5660	496/16649	2.860491e-10	1.143583e-07	9.486937e-08	RPS27L/INO80/SMARCB1/BACH1/RBBP8/IER3...	235
GO:0002480	GO:0002480	antigen processing and presentation of exog...	36/5660	45/16649	3.073965e-10	1.146999e-07	9.515270e-08	HLA-E/HLA-F/HLA-C/HLA-A/HLA-B/HLA-E/HL...	36

**NOTE:** Instead of saving just the results summary from the `ego` object, it might also be beneficial to save the object itself. The `save()` function enables you to save it as a `.rda` file, e.g. `save(ego, file="results/ego.rda")`. The complementary function to `save()` is the function `load()`, e.g. `ego <- load(file="results/ego.rda")`.

*This is a useful set of functions to know, since it enables one to preserve analyses at specific stages and reload them when needed.* More information about these functions can be found [here](#) & [here](#).

## Visualizing clusterProfiler results


clusterProfiler has a variety of options for viewing the over-represented GO terms. We will explore the dotplot, enrichment plot, and the category netplot.

The **dotplot** shows the number of genes associated with the first 50 terms (size) and the p-adjusted values for these terms (color). This plot displays the top 50 GO terms by gene ratio (# genes related to GO term / total number of sig genes), not p-adjusted value.

```
## Dotplot
dotplot(ego, showCategory=50)
```

To save the figure, click on the **Export** button in the RStudio **Plots** tab and **Save as PDF...**. In the pop-up window, change:

- Orientation: to Landscape
- PDF size to 8 x 14 to give a figure of appropriate size for the text labels



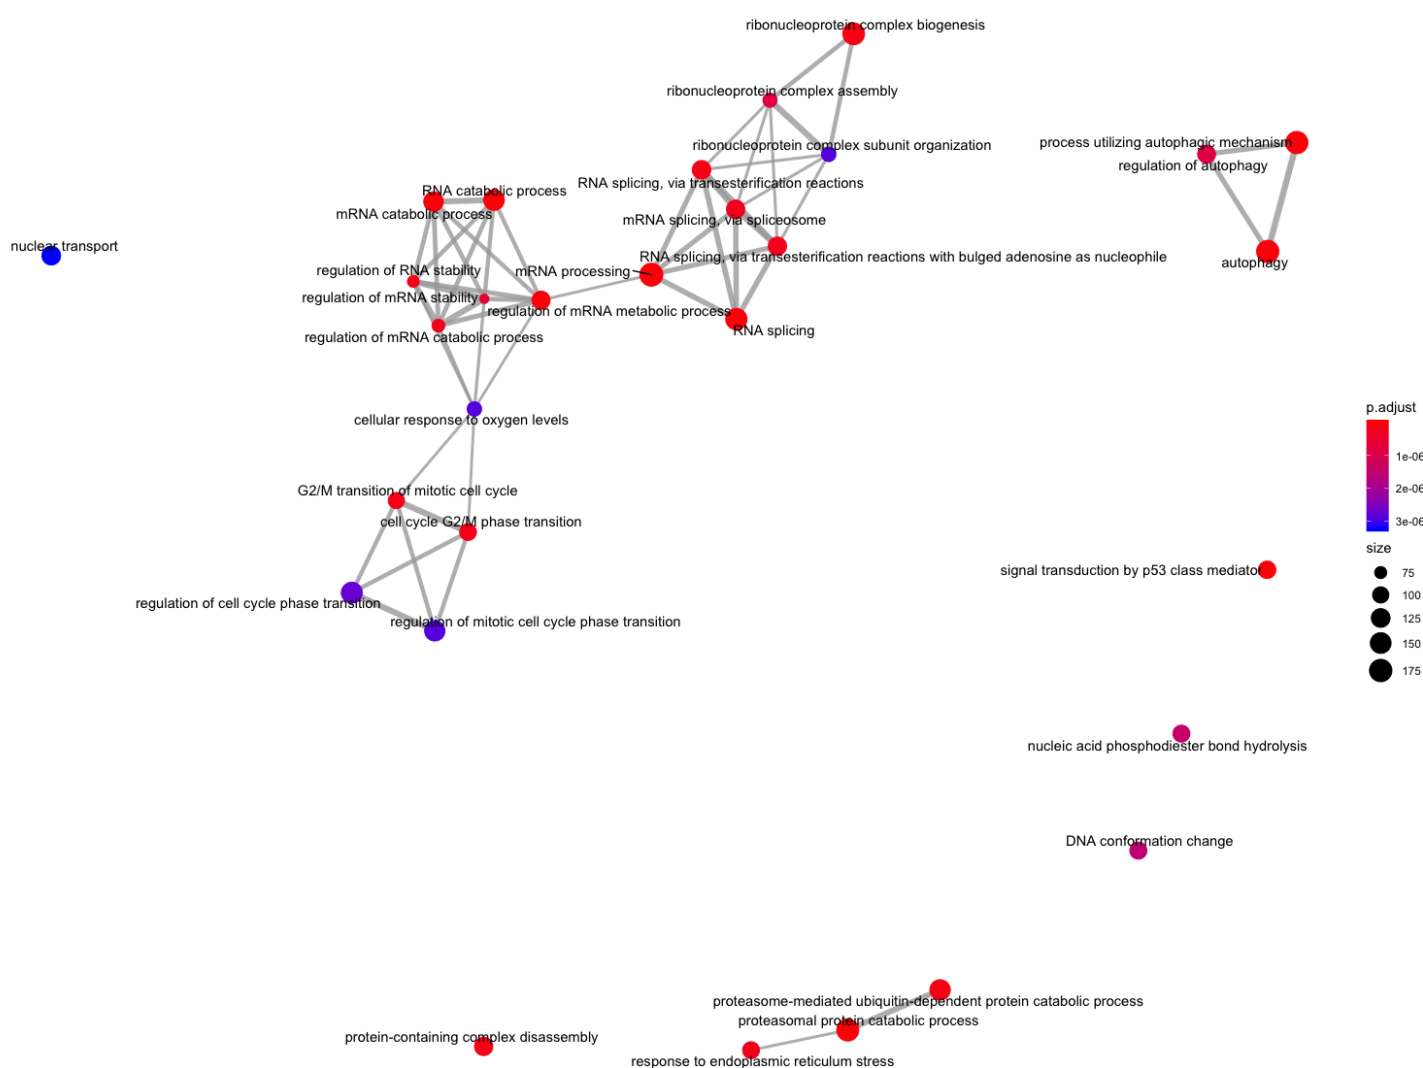


The next plot is the **enrichment GO plot**, which shows the relationship between the top 50 most significantly enriched GO terms (padj.), by grouping similar terms together. Before creating the plot, we will need to obtain the similarity between terms using the `pairwise_termsim()` function ([instructions for emapplot](#)). In the enrichment plot, the color represents the p-values relative to the other displayed terms (brighter red is more significant), and the size of the terms represents the number of genes that are significant from our list.

```
## Add similarity matrix to the termsim slot of enrichment result
ego <- enrichplot::pairwise_termsim(ego)

## Enrichmap clusters the 50 most significant (by padj) GO terms to visualize relationships
emapplot(ego, showCategory = 50)
```

To save the figure, click on the **Export** button in the RStudio **Plots** tab and **Save as PDF...**. In the pop-up window, change the **PDF** size to 12 x 14 to give a figure of appropriate size for the text labels.



Finally, the **category netplot** shows the relationships between the genes associated with the top five most significant GO terms and the fold changes of the significant genes associated with

these terms (color). The size of the GO terms reflects the pvalues of the terms, with the more significant terms being larger. This plot is particularly useful for hypothesis generation in identifying genes that may be important to several of the most affected processes.

**Note** - You may need to install the `ggnewscale` package using `install.packages("ggnewscale")` for the `cnetplot()` function to work.

```
## To color genes by log2 fold changes, we need to extract the log2 fold changes :
OE_foldchanges <- sigOE$log2FoldChange

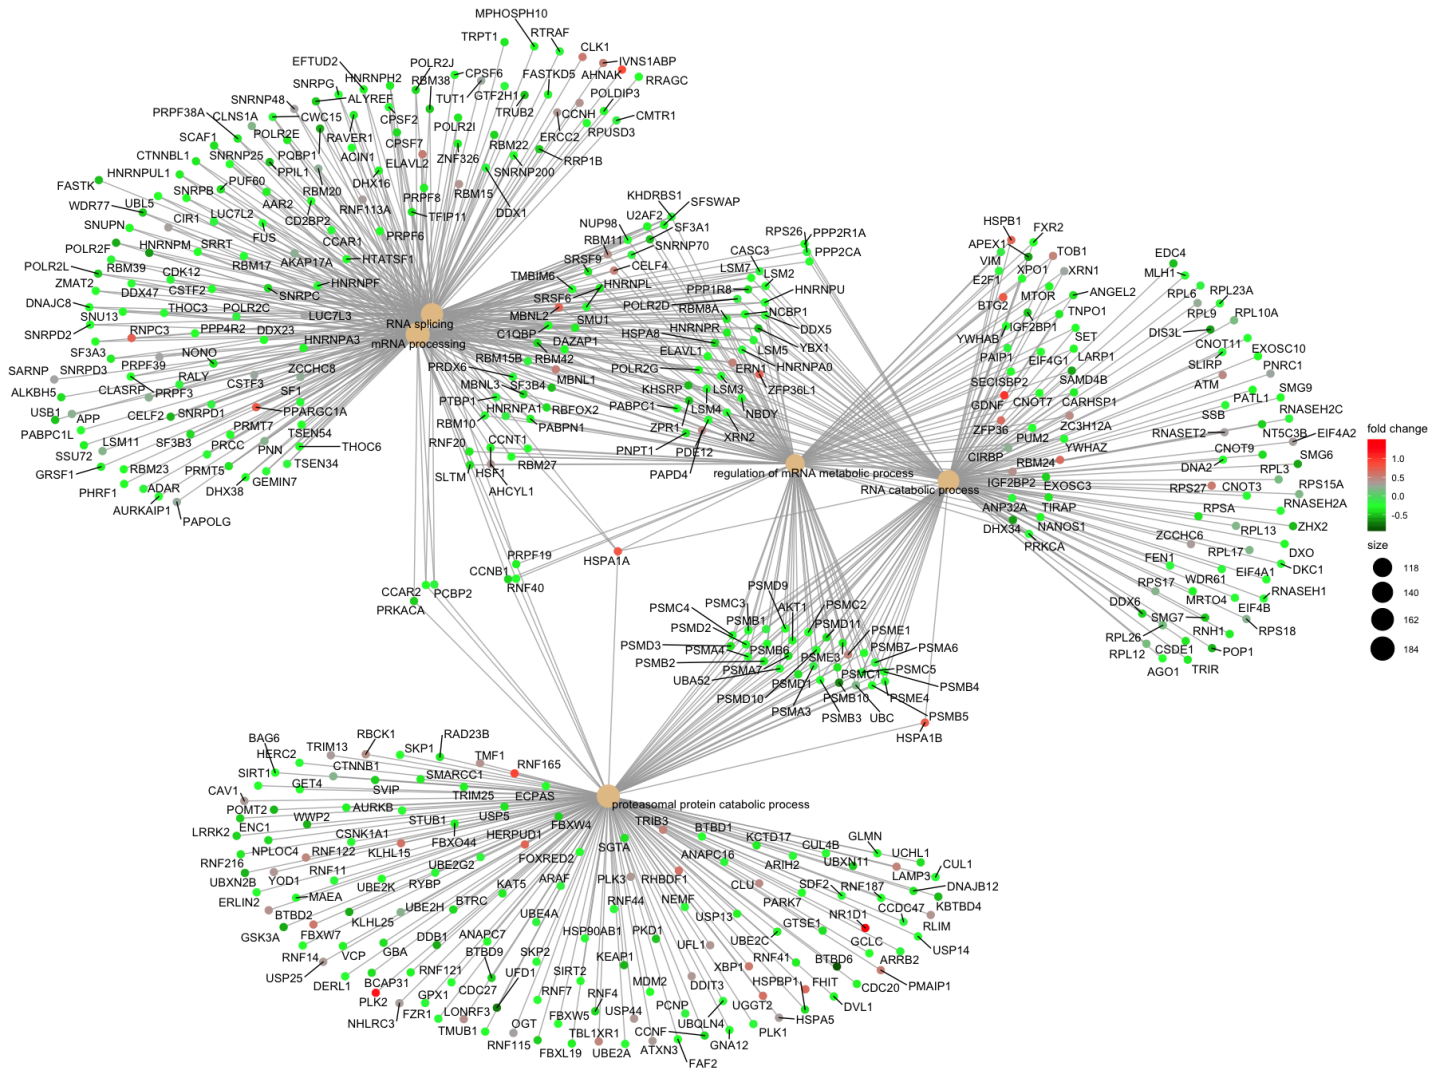
names(OE_foldchanges) <- sigOE$gene

## Cnetplot details the genes associated with one or more terms - by default gives
cnetplot(ego,
          categorySize="pvalue",
          showCategory = 5,
          foldChange=OE_foldchanges,
          vertex.label.font=6)

## If some of the high fold changes are getting drowned out due to a large range,
OE_foldchanges <- ifelse(OE_foldchanges > 2, 2, OE_foldchanges)
OE_foldchanges <- ifelse(OE_foldchanges < -2, -2, OE_foldchanges)

cnetplot(ego,
          categorySize="pvalue",
          showCategory = 5,
          foldChange=OE_foldchanges,
          vertex.label.font=6)
```

**Again, to save the figure,** click on the **Export** button in the RStudio **Plots** tab and **Save as PDF...** Change the **PDF** size to **12 x 14** to give a figure of appropriate size for the text labels.

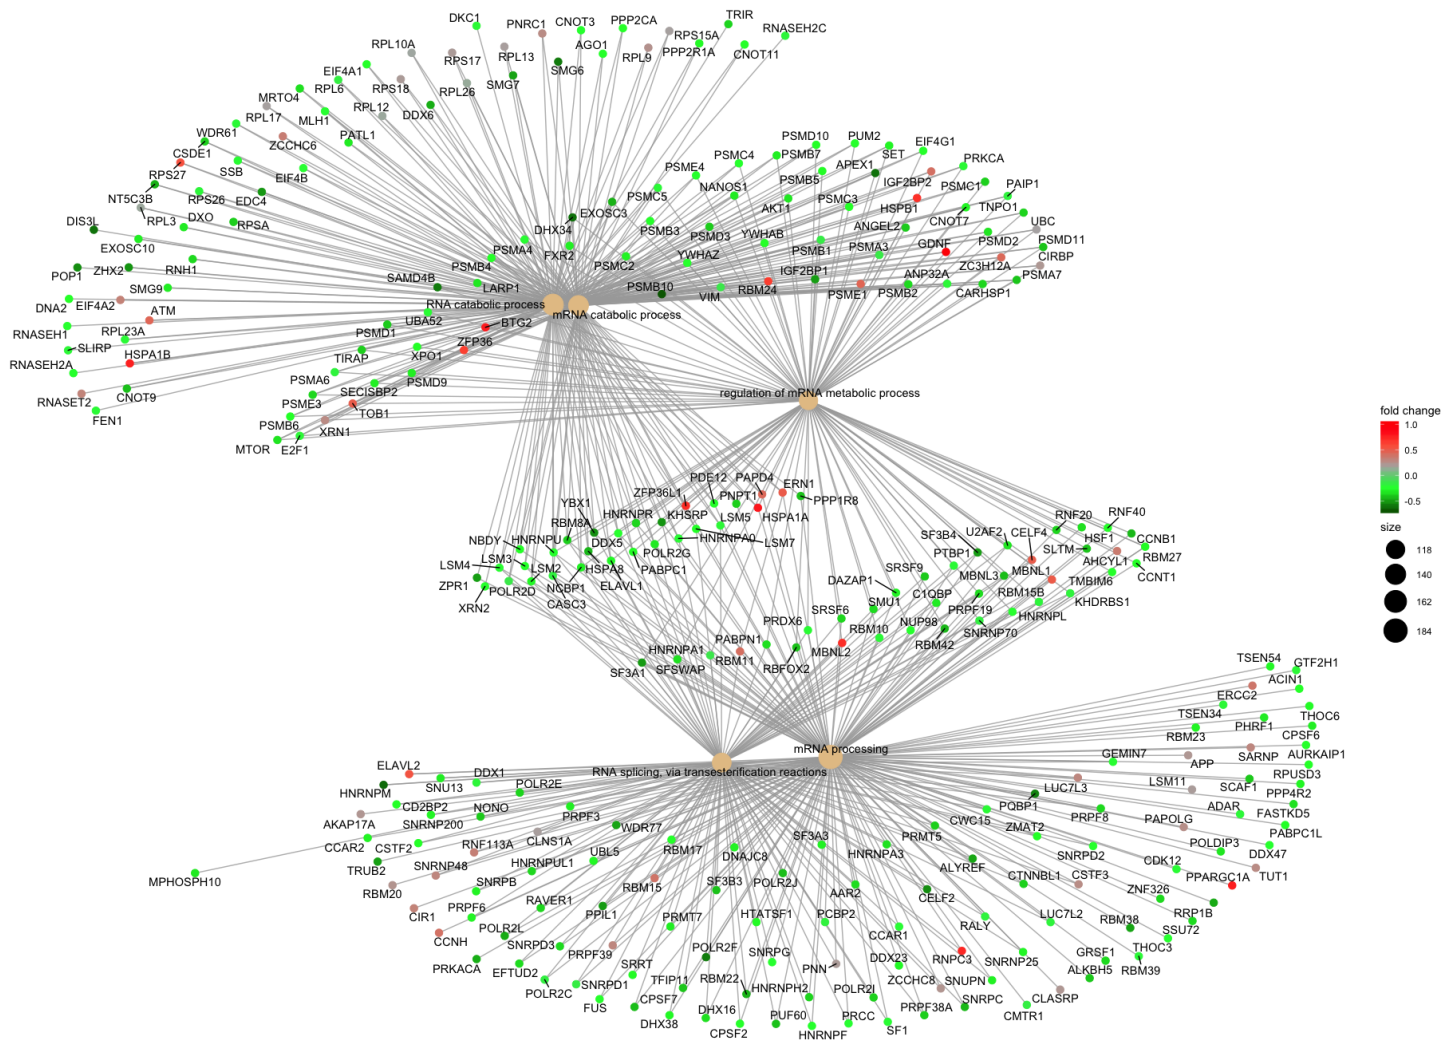


If you are interested in significant processes that are **not** among the top five, you can subset your `ego` dataset to only display these processes:

```
## Subsetting the ego results without overwriting original `ego` variable
ego2 <- ego

ego2@result <- ego@result[c(1,3,4,8,9),]

## Plotting terms of interest
cnetplot(ego2,
  categorySize="pvalue",
  foldChange=OE_foldchanges,
  showCategory = 5,
  vertex.label.font=6)
```



This lesson has been developed by members of the teaching team at the [Harvard Chan Bioinformatics Core \(HBC\)](#). These are open access materials distributed under the terms of the [Creative Commons Attribution license \(CC BY 4.0\)](#), which permits unrestricted use, distribution, and reproduction in any medium, provided the original author and source are credited.

**DGE\_workshop\_salmon\_online** is maintained by [hbctraining](#).

This page was generated by [GitHub Pages](#).